

Foundations of machine learning

Machine Learning in Finance · Session 4

Where three sessions of Python left you

Eleven stocks, ten years, ranked by how much they move:

```
ticker
NVDA    48.6
AAPL    28.5
DIS     27.8
XOM     27.8
JPM     27.4
MSFT    27.1
WMT     21.1
PG      18.5
JNJ     18.1
KO      17.8
SPY     17.6
```

Annualised volatility, in percent. Three lines of pandas.

A question that ranking cannot answer



Will tomorrow be an up day? Nothing you have written so far even attempts this. Every tool from Sessions 1 to 3 describes data that already happened.

Describing, and predicting

Describing

What did happen?

- Apple's volatility was 28.4% a year
- Nvidia moved most, Coca-Cola least
- The worst day was 16 March 2020

You can check every one of these against the data.

Predicting

What will happen?

- Tomorrow's return will be about +0.3%
- This client will default
- This filing is unusual

You cannot check any of them yet. That is the whole problem.

Today

A prediction needs four things: **a table of the right shape, a rule, a way to score the rule**, and **data the rule has never seen**. Today is those four things.

By the end of the session you can:

- put a problem into the **statistical learning framework**
- tell **supervised** from **unsupervised**, **regression** from **classification**
- handle **missing values, extreme values and columns on different scales**
- score a prediction with **MAE, MSE, RMSE, R^2** or a **confusion matrix**
- explain **overfitting** and the **bias-variance trade-off** to somebody else

This session is part maths

Sessions 1 to 3 were about writing code. This one is mostly about ideas, and ideas here are written in symbols.

Nothing today is new maths. It is the vectors, matrices and sums you already use in econometrics, written the way machine learning writes them, plus the occasional expectation.

We still write code, and there are live cells throughout. But **we do not fit a model today**. That starts in the next block, once the ideas are in place.

1 • The statistical learning framework

PART 1

Observations, features, targets, and the error nobody can remove.

It all starts with one table

features					target
date	ret_1d	ret_5d	vol_20d	volume	ret_next
2024-11-04	-0.0018	0.0121	0.0143	41.2M	0.0091
2024-11-05	0.0091	0.0087	0.0140	28.1M	0.0042
2024-11-06	0.0042	0.0102	0.0138	54.6M	-0.0033
2024-11-07	-0.0033	0.0064	0.0135	42.1M	-0.0121
2024-11-08	-0.0121	-0.0009	0.0137	38.3M	0.0107

n rows
= n obs.

Each row is one observation. Each feature column is something you know at that moment. The target column is what you want to predict.

The same objects, three vocabularies

You already know most of these words. Machine learning just uses different ones.

econometrics	machine learning	in this table
observation	sample, instance, example	one row, one date
independent variable, regressor	feature, predictor, input	<code>ret_1d</code> , <code>vol_20d</code>
dependent variable	target, label, response, output	<code>ret_next</code>
coefficient	weight, parameter	one number per feature
residual	error	what the rule got wrong
model	learner, estimator	the rule itself

Building that table, from prices

Python Code

```
prices = pd.read_csv("data/prices.csv", parse_dates=["date"])
wide = prices.pivot(index="date", columns="ticker", values="close")
rets = wide.pct_change()

aapl = pd.DataFrame({
    "ret_1d": rets["AAPL"],
    "ret_5d": rets["AAPL"].rolling(5).mean(),
    "vol_20d": rets["AAPL"].rolling(20).std(),
})
aapl["ret_next"] = rets["AAPL"].shift(-1)
aapl = aapl.dropna()
aapl.tail(4)
```

	ret_1d	ret_5d	vol_20d	ret_next
date				
2024-12-24	0.011478	0.003790	0.008828	0.003175
2024-12-26	0.003175	0.008710	0.008782	-0.013242
2024-12-27	-0.013242	0.004659	0.009615	-0.013264
2024-12-30	-0.013264	-0.001757	0.010262	-0.007058

`shift(-1)` pulls tomorrow's return back onto today's row. Every feature on a row is known **before** the target happens. That is the rule that makes the table honest.

n and p

The table has a standard notation, and it is the notation from Session 3.

$$\mathbf{X} \in \mathbb{R}^{n \times p}, \quad \mathbf{y} \in \mathbb{R}^n$$

- n is the number of **observations**, the rows
- p is the number of **features**, the columns of $\mathbf{X}$
- x_{ij} is the value of feature j for observation i
- y_i is the target for observation i

`aapl.shape` on the previous slide told you n and $p + 1$ directly. In NumPy $\mathbf{X}$ is a 2-dimensional array and $\mathbf{y}$ is a 1-dimensional one.


The assumption underneath everything

$$Y = \underbrace{f(X)}_{\text{signal}} + \underbrace{\varepsilon}_{\text{noise}}$$

There is some true relationship f between the features and the target, and whatever it does not account for is ε .


f is fixed and unknown. ε has mean zero and is independent of X . Everything in the rest of this course is an attempt to get at f .

The part that is really there




Invented data. With real data nobody can draw the black line, which is exactly why this picture is worth looking at now.

And the part that is not




ε is measurement error, features you did not collect, and genuine randomness, all at once. **In finance it is most of what you see.**

What you actually get



Thirty dots. No curve, no formula, and no way to check your answer against the truth. This is the situation you are always in.

So you estimate it



$\hat{f}$ is your estimate, built from the data you have. A **model** is a recipe for producing $\hat{f}$, and a **prediction** is $\hat{y} = \hat{f}(x)$.

Two reasons to want it

Prediction

You want the prediction to be close to the truth, and you do not care how.

- what will this loan lose?
- will this trade be profitable?
- is this transaction fraudulent?

The rule is allowed to be a black box.

Inference

You want to know **which features matter, and how**.

- does size explain returns after controlling for value?
- how much does one rating notch cost in spread?
- which factor drives this portfolio?

The rule has to be readable.

Econometrics leans hard on the right-hand column. Machine learning leans hard on the left. Same framework, different priority.

How wrong you can be, split in two

$$\mathbb{E} \left[(Y - \hat{Y})^2 \right] = \underbrace{\left[f(X) - \hat{f}(X) \right]^2}_{\text{reducible}} + \underbrace{\text{Var}(\varepsilon)}_{\text{irreducible}}$$


Reducible

The gap between your estimate and the truth. Better features, more data or a better model shrink it.

Irreducible

The noise itself. No model, no data and no amount of compute removes it. It is a property of the problem.

The floor, drawn



You can move the blue curve until it lies on top of the black one. **The red segments would still be there**, because they are the ε from earlier: the part of y that no feature in your table accounts for.

In finance the floor is high

0.56

R^2 of Apple's return on the market's,
same day

0.01

R^2 of Apple's return on **yesterday's**
return, in sample

-0.02

the same rule, on data it had not seen

Explaining today with today is easy. Explaining tomorrow with today is close to impossible, because $\text{Var}(\varepsilon)$ is enormous.

That last number is not a typo. We come back to it in Part 4.

You have already estimated an f

Session 3 ended on this line:

```
beta = np.linalg.inv(X.T @ X) @ X.T @ y
```

$$\hat{f}(x) = \hat{\beta}_0 + \hat{\beta}_1 x, \quad \hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Linear regression **is** a machine learning model: a rule with parameters, chosen to fit a sample. It is the simplest one, and it is still the one to beat.

Assume a shape, or do not

Parametric

Pick a form first, then find its numbers.

Linear regression fixes one coefficient per feature, plus an intercept, whatever the sample size.

Cheap, stable, easy to read. Wrong if the true shape is not that shape.


Non-parametric

Let the data choose the shape.

Nearest neighbours, trees, splines. The rule grows with the sample.

Flexible. Needs far more data, and is harder to explain.


The same data, both ways



Two shapes for the two coefficients you get on the left. On the right there are no coefficients at all: the rule is “average the five closest rows”.

The rows you learn from, and the rows you check on

A model has already seen the answers in the rows it was built from, and a flexible enough rule can simply memorise them. So a score computed on those rows is not evidence of anything.



Why you split is today's business, and it runs through the rest of the session. **How** to split well, and how to reuse the training rows without ever touching the test rows, is model selection and cross-validation in the next block.

Your turn 1

`rets["MSFT"]` is a Series of daily returns. Build a two-column table whose target is **the return five days ahead**.

Python Code

```
msft = pd.DataFrame({"ret_1d": rets["MSFT"]})

msft["ret_5d_ahead"] = ...
msft = msft.dropna()
msft.tail(3)
```

Your turn 1: answer

```
msft = pd.DataFrame({"ret_1d": rets["MSFT"]})  
  
msft["ret_5d_ahead"] = rets["MSFT"].shift(-5)  
msft = msft.dropna()  
msft.tail(3)
```


A negative shift moves the future backwards onto today's row. Shifting the wrong way gives you yesterday's return as a target, which you could predict perfectly and which would be worth nothing.

2 · What kind of problem is it?

PART 2

Supervised, unsupervised, and the one we are not covering.

One question decides almost everything



Supervised: you have the answers

For every row in the training data, somebody has already written down the target.

Credit

past loans, and whether each one defaulted

Trading

past days, and what the return turned out to be

Fraud


past transactions, and which were later confirmed as fraud

Ratings

past issuers, and the rating each was given


The target column is the expensive part. Somebody had to observe it, or label it by hand, and there is usually far less of it than you would like.

Regression: the target is a number



The answer can be any value on the vertical axis. **This is the OLS picture you already know**, seen as a prediction problem.

Classification: the target is a label



Now there are only two possible answers, and being “nearly right” is not a thing.

The same data, two different problems

Nothing changed in the data. **You chose** what the target would be, and that choice changes the model you can use, the metric you score it with, and what a mistake costs you.

As regression

predict the return, then act on its size

As classification

predict the direction, then act on the direction

Turning a number into a label always throws information away. A +0.02% day and a +6% day become the same “up”. Sometimes that is the right trade, because direction is easier to get right than size.

Making the label yourself

Python Code

```
labels = (rets["AAPL"] > 0).astype(int)

print(labels.tail(6).to_string())
print()
print("share of up days:", round(labels.mean(), 4))
```

date	
2024-12-23	1
2024-12-24	1
2024-12-26	1
2024-12-27	0
2024-12-30	0
2024-12-31	0

share of up days: 0.5322

Unsupervised: nobody wrote the answers down

There is no target column. You have $\mathbf{X}$ and nothing else, and the question becomes “what is the structure in here?”

Clustering

which of these stocks behave alike, without being told the sectors

Anomaly detection

which trades do not look like any of the others


Dimension reduction

500 correlated series, and three factors that carry most of it

Topic models

what are these ten thousand filings about

Eleven stocks, no target



Nobody labelled these. You can still see that the defensive names sit together in the bottom left, and Nvidia sits alone. **That is clustering, done by eye.**

The catch with unsupervised work

Supervised

There is a right answer, so there is a score.

You can say “this model is better than that one” and defend it with a number.

Unsupervised

There is no right answer, so there is no score.

Three clusters or five? Both are defensible. You are back to judgement.

This is why supervised learning dominates in practice, and why most of this course is about it.

Reinforcement learning

The third paradigm. No fixed table at all: an **agent** takes actions in an environment, gets rewards, and learns a policy that earns more of them.

In principle

optimal execution, hedging, market making, portfolio rebalancing

In practice

it needs an environment to practise in, and markets do not offer one

We do not cover reinforcement learning in this course. You should know the word, and know that it is a different setting: the data depends on what the agent decided to do.

Your turn **2**

Which is each of these?

1. Predict a bond's yield spread next month
2. Group 3,000 funds by how they behave, with no style labels
3. Decide whether a filing will be followed by a restatement
4. Compress 60 macro series into a handful of factors
5. Predict how many basis points a stock moves after an earnings surprise

Your turn 2: answer

1. Predict a spread → **regression**, the target is a number
2. Group funds → **unsupervised**, clustering, no target exists
3. Restatement or not → **classification**, the target is a label
4. Compress 60 series → **unsupervised**, dimension reduction
5. Basis points after a surprise → **regression**, the target is a number

Number 3 could be made a regression by predicting the probability instead. The line is not always sharp, and the decision is yours to defend.

3 • Real data is never that shape

PART 3

Missing values, extreme values, and columns
measured in different worlds.

Missing values, and where they come from

Calendars disagree

US and European exchanges keep different holidays, so a merged table has gaps

Trading stopped

halts, suspensions, and days with no quote

The last one is the dangerous one, and we come back to it.

The company was not there

an IPO in 2019 has no 2015 rows, and a delisted firm has no 2024 rows

Reporting is optional

a fund reports until the month it would rather not

Finding them


Python Code

```
calendar = pd.date_range(wide.index.min(), wide.index.max(), freq="D")
daily = wide.reindex(calendar)

print("rows before:", len(wide))
print("rows after :", len(daily))
print(daily.isna().sum().head(3).to_string())
```

```
rows before: 2516
rows after : 3652
ticker
AAPL      1136
DIS        1136
JNJ        1136
```

A picture of it



Two solid rows of weekends, and **ten single red squares**: New Year, Martin Luther King Day, Presidents' Day, Good Friday, and the rest of the market holidays. A picture of missingness usually tells you the cause immediately, which is what you need before you decide what to do about it.

Three ways out

Drop the rows

`df.dropna()`. Safe when the gaps are few and unrelated to what you are studying.

Fill them in

`fillna`, `ffill`, `interpolate`. You are inventing data, so be deliberate about it.


Drop the column

if a feature is 60% empty it is usually not worth rescuing.

Keep the fact

add a column that records “this was missing”. Sometimes that is the signal.

What each choice looks like



Grey dashed is the truth, red crosses are the invented values. Fourteen days removed from Apple's spring.

The accurate answer is the illegal one

\$10.86

mean error, fill with the column mean

\$2.76

mean error, forward fill

\$1.67

mean error, interpolate

Interpolation wins, because it draws a line between the price **before** the gap and the price **after** it.

Standing on 8 May you would not have known 9 May's price, and that is the price the interpolated value is drawn from. Interpolation quietly puts tomorrow's information into today's row, and every result built on it is optimistic. **Forward fill only ever looks backwards.**

Why it is missing decides what you may do

Missing for an unrelated reason

Thanksgiving. A vendor outage. A field nobody filled in.

Dropping those rows costs you sample size and nothing else.


Missing because of the value

The fund stopped reporting after the month it lost 40%. The firm delisted because it failed.

Drop those rows and your sample is only the survivors.

The second case has a name in finance: **survivorship bias**. It is the reason a backtest on today's index members looks so good.

Extreme values



Most of them are in March 2020, and they are not mistakes. **The best day was +12.0% and the worst was -12.9%, on the Friday and the Monday of the same weekend.**

An error, or the whole point?

Probably an error

A price of 0, or of 1,000,000. A return of +4,000% because of an unadjusted stock split. A date in 1901.

Fix it or drop it, and write down what you did.


Probably the point

March 2020. The 2008 quarter. The day the takeover was announced.

Delete these and you have built a model of a market where nothing ever happens.

In most fields the extremes are noise. **In finance they are usually the phenomenon.**
Risk lives in the tail, so removing the tail removes the subject.

Two rules of thumb, and their limits



Q_1 and Q_3 are the 25th and 75th percentiles, and the **interquartile range** $\text{IQR} = Q_3 - Q_1$ is the width of the middle half of the data.

$$\text{IQR rule: outside } [Q_1 - 1.5 \text{ IQR}, Q_3 + 1.5 \text{ IQR}] \quad \text{z-score rule: } \left| \frac{x - \bar{x}}{s} \right| > 3$$

Both rules flag far too much

138

days flagged by the IQR rule, 5.5% of the sample

36

days beyond three standard deviations


7

days a normal distribution says you should get

The rules assume a bell curve. Returns are not one, so the rules mark hundreds of perfectly real days as errors.

There is also a circular problem: one enormous day inflates s , which then makes that day look less extreme.

The tails are the whole story



The bars sit far above the red line at both ends. **Days that “cannot happen” happen several times a decade**, and they are the ones that decide whether a strategy survives.

So what do you actually do

Check the source first

most true errors are visible in one plot or one
`.describe()`

Score with a robust metric

MAE cares far less about one bad day than MSE does. Part
4.

Do not delete by default


a rule that flags 5% of your days is not finding errors, it is
finding the market

Or transform

a log squeezes the long right tail of prices, volumes and
market caps

Whatever you choose, **write it down**. “We dropped returns beyond $\pm 20\%$, three observations” is a sentence that belongs in the report.

Levels, and changes



The left panel wanders off and never comes back, so its mean means nothing. The right panel hovers around a fixed level, so its mean and standard deviation describe something real. **Model the changes, not the levels.**

Why logs

$$r_t^{\log} = \log\left(\frac{P_t}{P_{t-1}}\right) = \log P_t - \log P_{t-1}$$

They add up

a year's log return is the sum of its days. Simple returns have to be compounded.

They tame right skew

market cap, volume and price are all long-tailed, and a log makes them workable


They are symmetric

-50% then +100% gets you back to where you started, and in logs those cancel exactly

They are nearly the same

for small daily moves, $\log(1 + r)$ and r agree to three decimals

Columns from different worlds



The volume column varies about **four billion times** more than the return column. Any method that measures distance, or penalises large coefficients, will hear only the volume.

Standardising

$$z_i = \frac{x_i - \bar{x}}{s}$$

Python Code

```
vol20 = rets["AAPL"].rolling(20).std().dropna()
z = (vol20 - vol20.mean()) / vol20.std()

print("original : mean %.5f  sd %.5f" % (vol20.mean(), vol20.std()))
print("z-scored : mean %.5f  sd %.5f" % (z.mean(), z.std()))
```

```
original : mean 0.01627  sd 0.00764
z-scored  : mean -0.00000  sd 1.00000
```

Every column now has mean 0 and standard deviation 1, so a coefficient of 0.4 means the same thing in every column. OLS does not need this. Nearest neighbours, penalised regression and neural networks all do.

Columns that are not numbers

Python Code

```
firms = pd.DataFrame({
    "ticker": ["AAPL", "KO", "JPM", "XOM"],
    "sector": ["Tech", "Staples", "Financials", "Energy"],
})

pd.get_dummies(firms, columns=["sector"], dtype=int)
```

	ticker	sector_Energy	sector_Financials	sector_Staples	sector_Tech
0	AAPL	0	0	0	1
1	KO	0	0	1	0
2	JPM	0	1	0	0
3	XOM	1	0	0	0

One column of text becomes one 0/1 column per category. This is **one-hot encoding**, and it is how a sector, a rating or a country gets into **X**.

The rule that governs all of it

Every number you use to clean the data, the mean you filled with, the standard deviation you divided by, the quantiles you flagged outliers with, **must be computed on the training rows alone.**

Those are the rows from Part 1 that the model is allowed to learn from. Compute the mean over the whole sample instead and it already contains the period you were going to test on, so your test score flatters you and the model disappoints in production.

It is the same look-ahead problem as interpolating across a gap, and it is the single most common mistake in student projects and in published finance papers alike.

Your turn 3

Apply the IQR rule to Microsoft's returns, and count how many days it flags.

Python Code

```
msft_ret = rets["MSFT"].dropna()
q1, q3 = msft_ret.quantile([0.25, 0.75])

iqr = ...
flagged = ...

print("days flagged:", flagged)
print("out of      :", len(msft_ret))
```

Your turn 3: answer

```
msft_ret = rets["MSFT"].dropna()
q1, q3 = msft_ret.quantile([0.25, 0.75])

iqr = q3 - q1
flagged = ((msft_ret < q1 - 1.5 * iqr) | (msft_ret > q3 + 1.5 * iqr)).sum()

print("days flagged:", flagged)
print("out of      :", len(msft_ret))
```

134 days of 2,515, or 5.3%. Almost exactly what the rule flagged for Apple, and far too many to be errors. The fences sit at -3.2% and +3.5%, so the rule is calling an ordinary bad Tuesday an outlier.

4 • Was the prediction any good?

PART 4

Turning a column of mistakes into one number you can defend.

A rule, and what it predicted

Session 3 fitted Apple's return against the market's and got two numbers. That is a model:

$$\hat{y}_t = 0.00043 + 1.2073 x_t$$

Python Code

```
market = rets["SPY"].dropna()
actual = rets["AAPL"].dropna()


predicted = 0.00043 + 1.2073 * market
resid = actual - predicted

pd.DataFrame({"actual": actual, "predicted": predicted, "resid": resid}).tail(3)
```

	actual	predicted	resid
date			
2024-12-27	-0.013242	-0.012279	-0.000964

The residual

$$e_i = y_i - \hat{y}_i$$



Why you cannot just add them up

Python Code

```
print("sum of residuals :", round(resid.sum(), 4))  
print("mean residual    :", round(resid.mean(), 8))  
print("mean size of one  :", round(resid.abs().mean(), 8))
```

```
sum of residuals    : -0.0096  
mean residual       : -3.81e-06  
mean size of one    : 0.00831654
```

The average residual is four millionths. The average residual **size** is more than two thousand times bigger. Least squares picked the line that makes the positives and the negatives cancel, so **the sum tells you nothing at all about how wrong the model was.**

You need to remove the signs first. There are two standard ways to do that.

Mean absolute error

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

The average size of a mistake, in the units of the target.

For this model, 0.83%. On a typical day the rule is wrong by about eight tenths of a percentage point, and that sentence is understandable by anybody.

Mean squared error, and its root

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad \text{RMSE} = \sqrt{\text{MSE}}$$

Squaring removes the sign too, but it also **punishes a big miss much harder than several small ones.**

MSE is in squared units, which nobody can interpret: 0.000142 squared returns. Taking the root puts it back into percent, which is why RMSE is the one you report.

All three, on the real residuals

Python Code


```
mae = resid.abs().mean()
mse = (resid ** 2).mean()
rmse = np.sqrt(mse)

print("MAE  %.5f  (%.2f%%)" % (mae, mae * 100))
print("MSE  %.6f" % mse)
print("RMSE %.5f  (%.2f%%)" % (rmse, rmse * 100))
```

```
MAE  0.00832  (0.83%)
MSE  0.000142
RMSE 0.01191  (1.19%)
```


RMSE is always at least as large as MAE. The gap between them tells you how uneven the errors are.

Why both exist



Being wrong by 4 costs **four times** as much as being wrong by 1 under MAE, and **sixteen times** as much under MSE.

One terrible forecast



The damage, in numbers

×1.8

MAE, 0.75% to 1.38%

×5.1

RMSE, 0.99% to 5.09%

×26

MSE

One day in sixty. **MAE says the model got a little worse. MSE says it fell apart.**

Neither is lying. They answer different questions: “how wrong is this model on a normal day” and “how badly can this model hurt me”.

A warning about MAPE

$$\text{MAPE} = \frac{100}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right|$$

Percentage error is attractive because it has no units, and it is common in sales and demand forecasting.

Never use it on returns. A day where the true return is 0.0001 makes the denominator almost zero, so one ordinary day can contribute thousands of percent to the average.


R^2 , and what it compares against

$$R^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2} = 1 - \frac{\text{RSS}}{\text{TSS}}$$

MAE and RMSE tell you how wrong you were. R^2 tells you how wrong you were **compared with the laziest possible model**: always predict the average.


$R^2 = 0$ means you matched that lazy model. $R^2 = 1$ means you fit every point exactly.

The two sums, drawn



R^2 is one minus the ratio of the red lengths to the amber ones, squared and added up.

R^2 on new data can be negative



Out of sample there is nothing forcing R^2 to be positive. **Negative means you would have done better by always predicting the average.**

And a tiny R^2 can still be worth a lot

An R^2 of 0.01 on next-day returns would be a serious result. It means 99% of what happens is noise to you, and the remaining 1%, applied every day across hundreds of names, is a business.

This is why finance is unlike most machine learning applications. A vision model with 1% explained variance is broken. A return model with 1% is a fund. Judge the number against the problem, never against 1.

Now change the target

Everything since the start of this part assumed the target was a **number**, so “how far off were you” had an answer. Change the target to a **label** and that question stops making sense: what is “up” minus “down”?

Regression, so far

predict Apple's return from the market's return

score it by the size of the residuals: MAE, RMSE, R^2

Classification, from here

predict whether Apple *rose*, from whether the market rose

there is no residual, so score it by counting right and wrong

Same eleven stocks, same 2,515 days, same market feature. Only the target column changed, exactly as in Part 2.

Four things can happen

You said up, and it went up

a true positive

You said up, and it went down

a false positive, a false alarm

You said down, and it went up

a false negative, a miss

You said down, and it went down

a true negative

Every classification metric in use is built out of those four counts, and nothing else.

The confusion matrix

		WHAT THE MODEL SAID	
WHAT HAPPENED		up	down
	it went up	1042 true positives	297 false negatives (misses)
	it went down	331 false positives (false alarms)	845 true negatives

Predicting “Apple rose today” from “the market rose today”, over 2,515 days.

Counting them yourself

Python Code

```
said_up = market > 0
went_up = actual > 0

tp = (said_up & went_up).sum()
fp = (said_up & ~went_up).sum()
fn = (~said_up & went_up).sum()
tn = (~said_up & ~went_up).sum()

print("TP %d   FP %d   FN %d   TN %d" % (tp, fp, fn, tn))
print("accuracy: %.4f" % ((tp + tn) / (tp + fp + fn + tn)))
```

```
TP 1042   FP 331   FN 297   TN 845
accuracy: 0.7503
```

Four boolean masks and four sums. Nothing here that Session 3 did not already give you.

Accuracy

$$\text{accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

75.0% of days called correctly. It is the first number everybody computes, and on its own it is close to useless.

Where accuracy falls apart

Now predict something rare: a day when Apple falls more than 5%.

22

such days in ten years

99.1%

accuracy of a model that always says "no crash"

0

crashes it would ever warn you about

A model with no skill whatsoever beats almost anything you will build, measured by accuracy. **When one class is rare, accuracy measures the class balance, not the model.**

Precision and recall

$$\text{precision} = \frac{TP}{TP + FP} \quad \text{recall} = \frac{TP}{TP + FN}$$

Precision: 75.9%

Of the days I called up, how many really were?

Low precision means **false alarms**: you traded on a signal that was not there.


Recall: 77.8%

Of the days that really were up, how many did I catch?

Low recall means **misses**: the move happened and you were not in it.

Neither uses the true negatives, which is exactly why they survive an imbalanced problem when accuracy does not.

You choose where to sit



Demand a score above +1% before calling a day up, and precision climbs to **91%** while recall collapses to **33%**.


One number, if you must have one

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

The harmonic mean, which is close to the smaller of the two. **Our model scores 0.768.**

A plain average would let 100% precision and 1% recall look respectable. The harmonic mean does not: it lands at 2%.

The whole trade-off, in one curve



Every threshold is one point on that curve. **AUC** is the area under it: 0.5 is a coin flip, 1.0 is perfect.

Which metric, and who decides

the question	the metric	why
how far off is a typical forecast	MAE	in the units of the target, robust
how badly can it hurt me	RMSE	large errors dominate, which is the point
is it better than predicting the mean	R^2	scale free, comparable across targets
the target is a label, classes balanced	accuracy	simple and honest here
the target is rare	precision and recall	accuracy is meaningless
I need to compare across thresholds	ROC AUC	threshold free

The metric is a business decision, not a technical one. What does a false alarm cost, and what does a miss cost?

Your turn 4

A “crash warning” model says the day will be bad whenever the market itself is down more than 1%. Compute its **recall** on Apple’s 5% down days.

Python Code

```
crash = actual < -0.05
warned = market < -0.01

hits = ...
recall = ...

print("crash days:", crash.sum())
print("recall    :", recall)
```

Your turn 4: answer

```
crash = actual < -0.05
warned = market < -0.01

hits = (warned & crash).sum()
recall = hits / crash.sum()

print("crash days:", crash.sum())
print("recall      :", recall)
```


It catches 19 of the 22 crashes, a recall of 0.864. Before you celebrate, ask the other question: it warns on 278 days in total, so its precision is 0.068. Thirteen false alarms for every real one.

5 · Overfitting, and the trade-off

PART 5

Why the model that fits best is usually the one to throw away.

A model that is never wrong



The same invented data from Part 1, and a curve that passes through almost every point. **By every metric in Part 4, this is an excellent model.**


So: is that the model you want?

Before the next slide, decide what you think.

- What does this curve say will happen at $x = 0.55$?
- If I collected thirty new observations, would it still fit them?
- Which parts of that curve are the relationship, and which are the noise?


There is no trick. The answer is the whole of the rest of this session.

Three fits of the same thirty points



Grey dashed is the truth. **Training error falls every time you add flexibility, and it will keep falling all the way to zero.**

The same three, on data they never saw



Four hundred new points. The middle model is now the best by a distance, and the flexible one is **thirty times worse** than it was on its own training data.

Two ways to be wrong

Underfitting

The model is too rigid to represent the relationship.

Bad on the training data, bad on new data, and **bad in the same way on both.**

The straight line above.

Overfitting


The model has learned the noise as if it were signal.

Excellent on the training data, poor on new data. **The gap is the symptom.**

The degree-15 curve above.

You cannot tell them apart from the training score. That is the entire reason for everything on the next few slides.

The picture worth remembering




Training error falls forever. Test error falls, bottoms out, and then climbs. Every model selection procedure in this course is a way of finding that bottom without cheating.

Turn the dial yourself

Python Code

```
degree = 3 # try 1, then 4, then 15
coef = np.polyfit(x_train, y_train, degree)

fig, ax = plt.subplots(figsize=(8.0, 1.5))
ax.scatter(x_train, y_train, s=26, color="#b8860b", zorder=3)
ax.plot(grid, np.polyval(coef, grid), color="#1c5cab", linewidth=2.2)
ax.set_ylim(-2.4, 2.4)
plt.show()
```



Which is why you held rows back

Training error kept improving all the way to degree 15. Test error stopped improving at 5. **Only one of those two numbers was available while the curve was being chosen.**

The rows you set aside in Part 1 are what turn “my model fits well” into evidence. Everything after today, cross-validation included, is a way of estimating that test error better without spending the test set.

Splitting the table

Python Code


```
train = aapl.loc[:"2022-12-31"]
test = aapl.loc["2023-01-01":]

print("train rows:", len(train), "from", train.index.min().date(),
      "to", train.index.max().date())
print("test rows :", len(test), "from", test.index.min().date(),
      "to", test.index.max().date())
```

```
train rows: 1994 from 2015-02-02 to 2022-12-30
test rows : 501 from 2023-01-03 to 2024-12-30
```

`aapl` is the feature table from Part 1. Splitting it is one line, because a DataFrame index knows about dates.

Time series break the usual rule



In the top panel the model trains on Wednesday and Friday to predict Thursday. **It knows the future.** Any backtest built that way will look wonderful and lose money.

The many faces of look-ahead bias

Shuffling a time series

the split has to be a date, not a random draw

Interpolating across a gap

Part 3: the fill used the price after the gap

Restated fundamentals

the numbers in the database today are not the ones published then

Scaling on the full sample

the mean you subtracted contains the test period

Choosing features after peeking

you tried forty ideas and kept the one the test set liked

Today's index members

the failures dropped out, so the sample is only the survivors

Where the error comes from

Test error is not one thing. Split it into two, and everything about model choice becomes clearer.

Bias

Error from the model being too simple to represent the truth.


It would still be there if you had a million rows.

Variance

Error from the model being too sensitive to the particular sample you drew.


Collect a different sample, get a different model.

The same model, twenty different samples



Left: every line is wrong in the same way, and they all miss the curve. Right: the lines are right on average and **wildly different from each other**.

Both, on one target



Each dot is the model you would have got from one sample. **Bias is how far the cluster sits from the centre. Variance is how spread out it is.**


The decomposition

$$\mathbb{E} \left[\left(y_0 - \hat{f}(x_0) \right)^2 \right] = \underbrace{\text{Var} \left(\hat{f}(x_0) \right)}_{\text{variance}} + \underbrace{\left[\text{Bias} \left(\hat{f}(x_0) \right) \right]^2}_{\text{bias}^2} + \underbrace{\text{Var}(\varepsilon)}_{\text{irreducible}}$$

Three non-negative terms. Expected test error can never fall below the third one, and **making one of the first two smaller usually makes the other larger.**

This is a statement about expected error over all possible training samples, not about the one sample you happen to have.

The trade-off, drawn



Bias falls, variance rises, and their sum has a minimum somewhere in between. **The red curve is the same U as before, now explained.**

The dial here is the number of neighbours rather than a polynomial degree, since it separates the two pieces more cleanly. Flexibility is the general idea; the particular knob does not matter.

What you actually do about it

Get more data

variance falls as n grows, and nothing else is as reliable

Simplify the model

a linear model is hard to overfit, which is why it is still the benchmark

Cross-validate

reuse the training data to estimate test error, without touching the test set

Use fewer features

every extra column is another way to fit noise

Penalise complexity

ridge and lasso deliberately add bias to remove more variance

Hold something back

and do not look at it until you have finished deciding

The last three are the subject of the next block. You now know what problem they are solving.

There is no best model

There is no model that is best for every problem. Whether flexibility helps depends on the shape of the true f and on how much noise sits on top of it, and you do not get to see either one.

So the job is not to find the best model. **It is to make the trade-off deliberately, and to be honest about how you measured it.**

What stays with you

One table

Rows are observations, columns are features, one column is the target.

Supervised or not

A target column decides it. Number means regression, label means classification.

Nothing from the future

Not in a fill, not in a scaler, not in a split.

Training error always falls

Which is exactly why it cannot be used to choose a model.

$$Y = f(X) + \varepsilon$$

You estimate f . You never remove ε .

Cleaning is the work

Missing, extreme, wrong scale, not numeric. In that order, most days.

Pick the metric first

MAE, RMSE and R^2 answer different questions. So do precision and recall.

Bias against variance

Too simple misses the shape. Too flexible learns the noise.

Where this goes next

now	the framework, the metrics, the trade-off	no models fitted
next	model selection and cross-validation	finding the bottom of the U
then	regression and classification methods, trees, ensembles	scikit-learn

Everything from here is a different way of building $\hat{f}$. The questions you ask of it are the ones from today, and they do not change.

The course case, Part 4 **CASE**

Four sessions ago the case asked which of two stocks was riskier. It has grown every session since:

Part 1

two stocks, a first look, arithmetic by hand

Part 3

all eleven stocks, pandas, and figures

Part 2

proper volatility, with loops and a function

Part 4

the same data, reframed as a prediction problem

Case Part 4: this week

You will take the table you built in Part 3 and turn it into a learning problem:

- choose a target, and build the features that are known before it
- find the missing values and the extremes, and decide what to do about each
- split by date, and say why the split has to be a date
- choose the metric you would judge a model by, and defend the choice
- state the prediction question precisely enough for somebody to answer it

You still do not fit anything. Part 4 finishes the setup, so that the first model you ever fit is one you already know how to judge.

Before the next session

- **Exercises:** the statistical learning framework, data problems, metrics, and over- and underfitting. Longer than usual, and worth the time.
- **Case Part 4:** the risk report becomes a prediction problem.

Stuck? Ask a friend, ask an AI for a hint rather than an answer, or email me at jobo@econ.au.dk.